

Progress Presentation

June 15th - June 27th


Agenda

- ❖ description of final (?) algorithm
- ❖ plots for choosing appropriate parameters
- ❖ the issue of “mismatch” between predicted and actual action
- ❖ other progress
- ❖ discussion of schedule




Final (?) algorithm

Planning



Final (?) algorithm

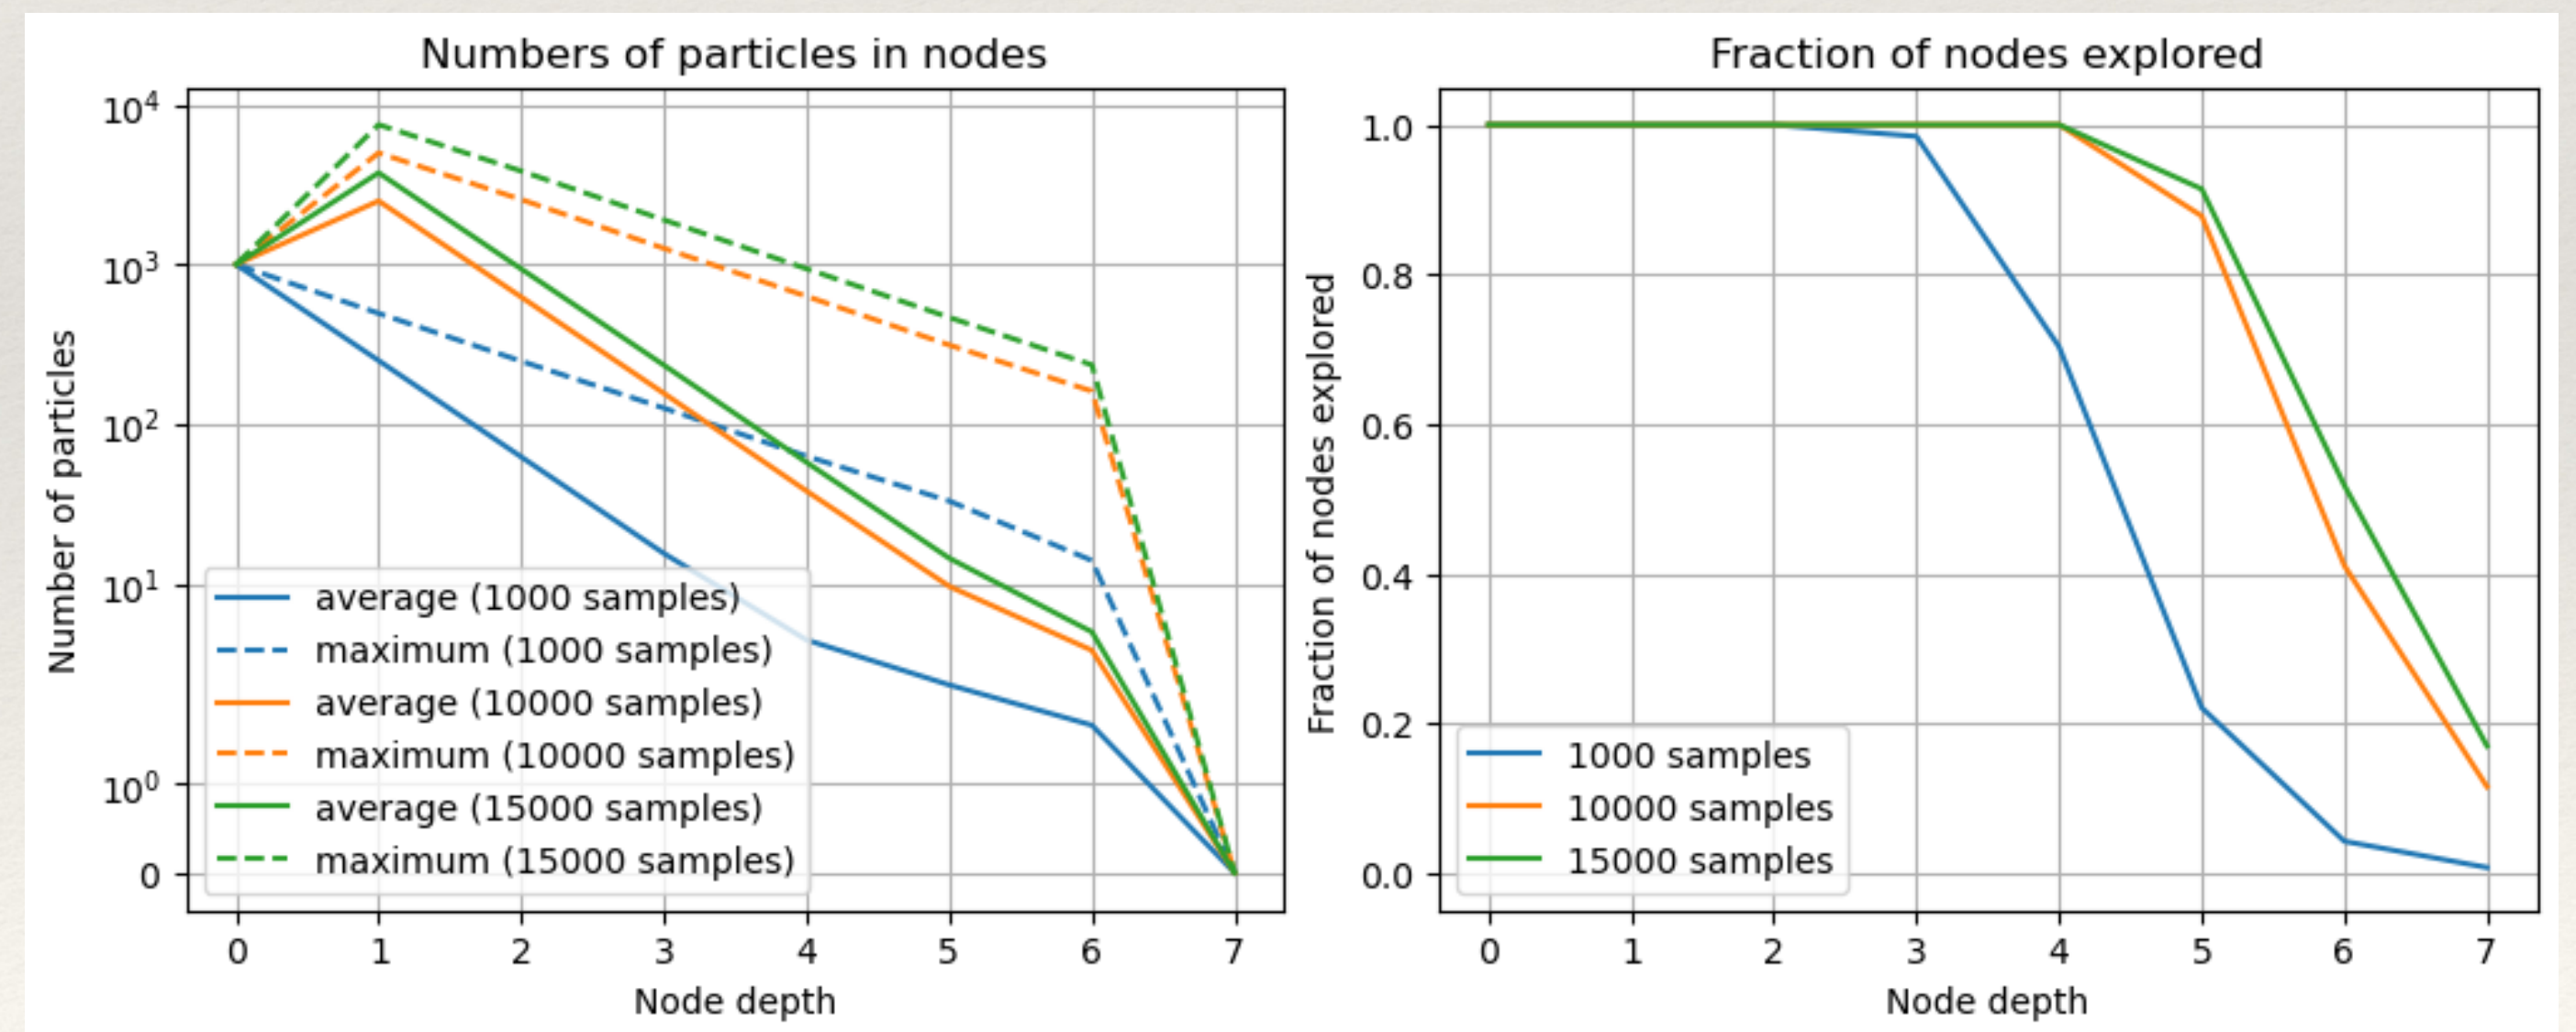
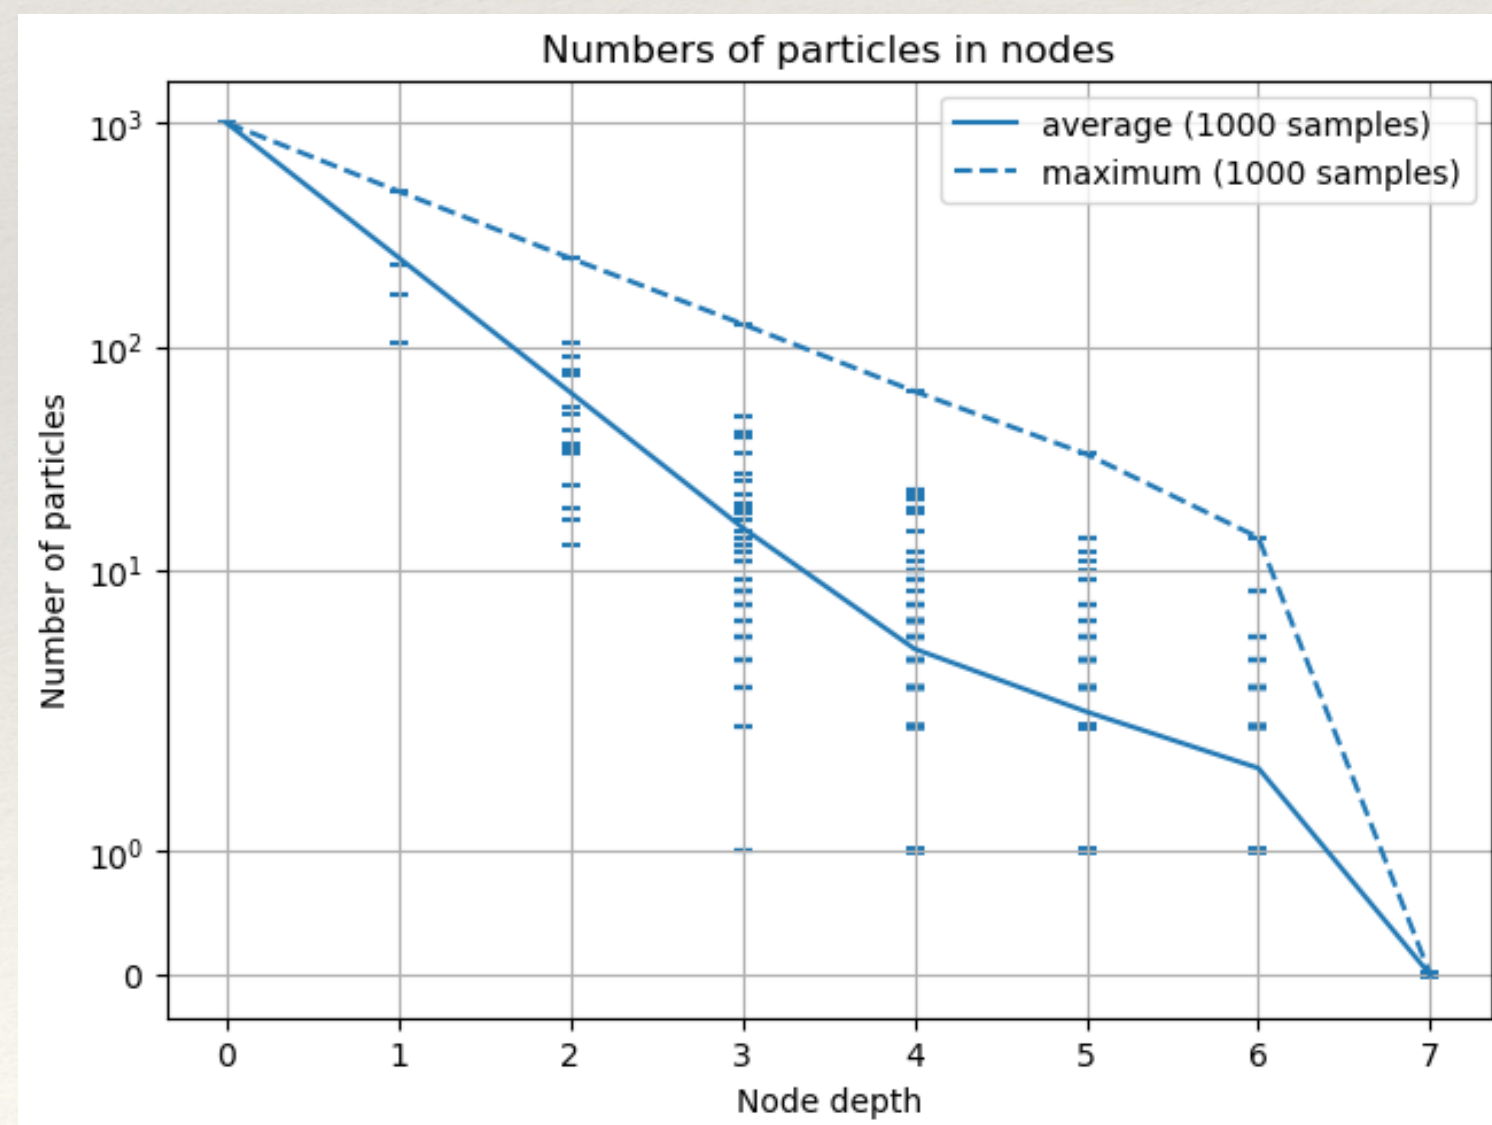
Belief Update



Choosing Solver Parameters

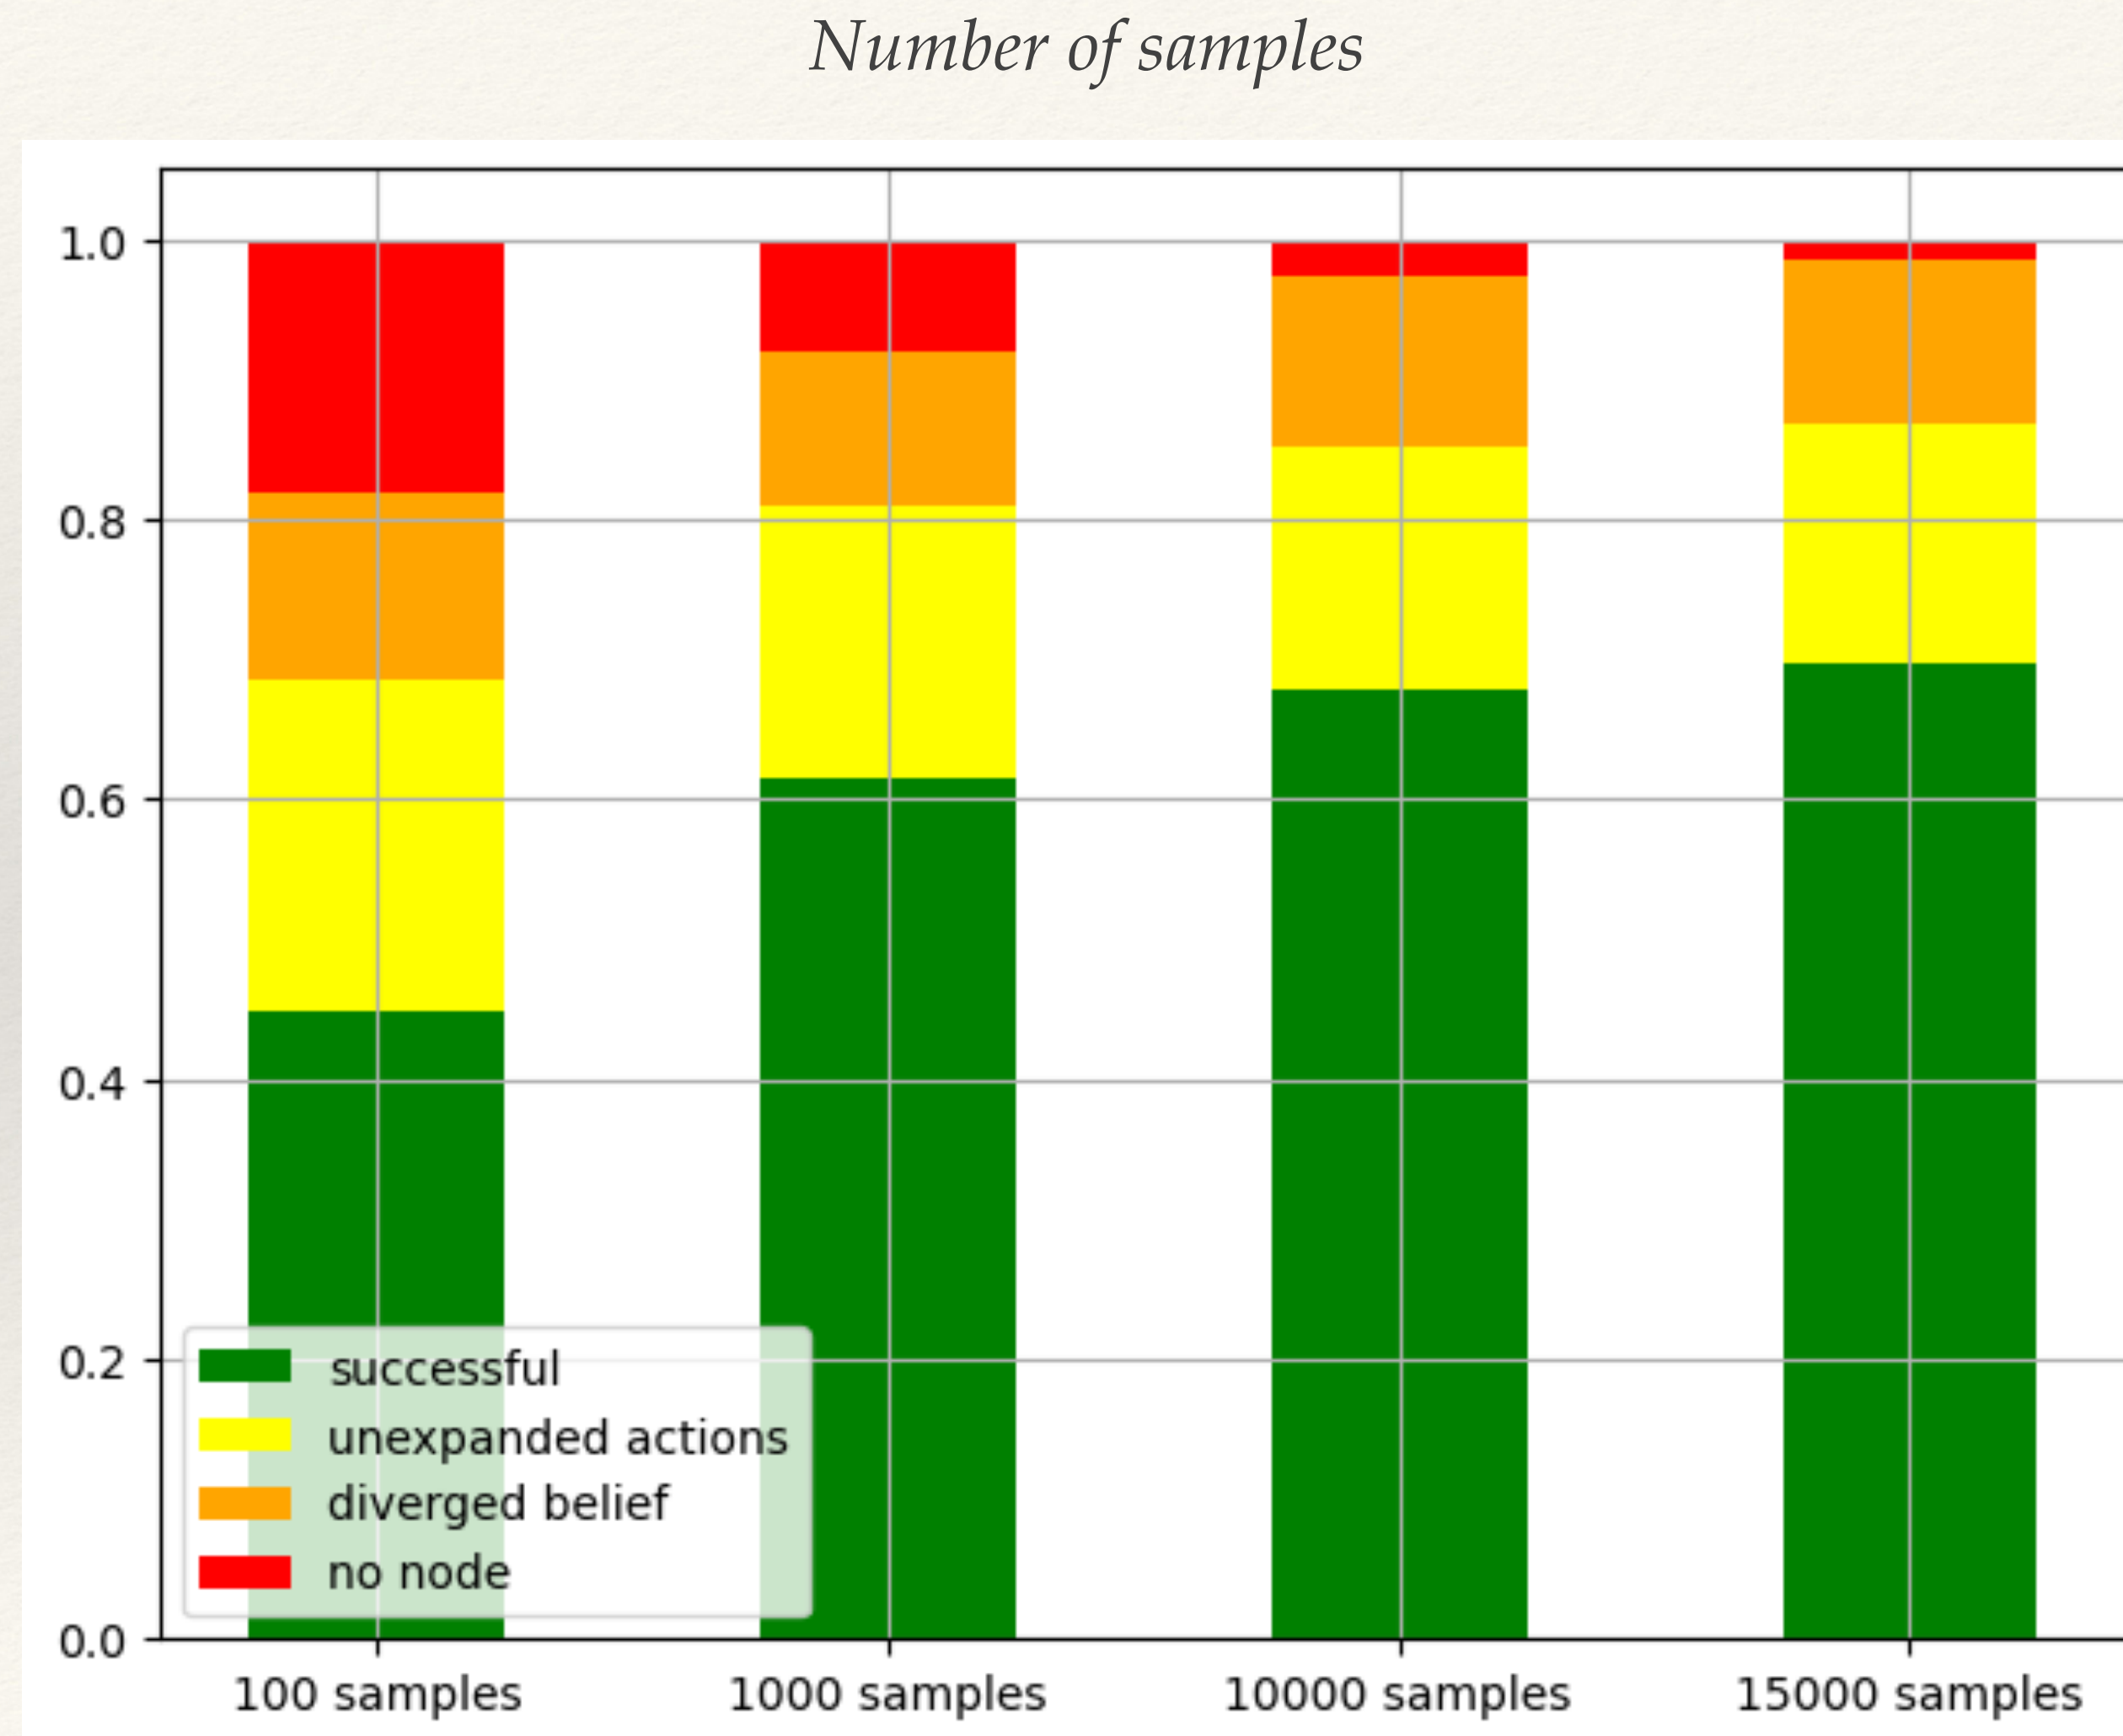
- one simulation
- two agents
- exploration coefficient 0.5
- reasoning level 0
- only first time step
- results from one agent's tree

Number of samples



Choosing Solver Parameters


- one simulation
- two agents
- exploration coefficient 0.5
- reasoning level 1
- only first time step
- results from one agent's tree



Choosing Solver Parameters

- one simulation
- 10k samples
- exploration coefficient 0.5
- reasoning level 0
- only first time step
- results from one agent's tree


Influence of number of agents



Choosing Solver Parameters

- one simulation
- two agents
- 10k samples
- reasoning level 1
- only first time step
- results from one agent's tree (plots 1 and 2 from level 0 tree, plot 3 from level 1 tree)


Exploration coefficient



Choosing Solver Parameters

- one simulation
- 10k samples
- reasoning level 0
- only first time step
- results from one agent's tree

Discount horizon depth



The Mismatch Issue

- ❖ what should an agent do when it predicted some action but observed another?
 - ❖ this can cause the particle filter to diverge in the lower tree when simulating the tree above
- ❖ I-POMDP does not seem to have an answer: interactive states $is_{1,l} = (s, \theta_{2,l-1}, \dots, \theta_{n,l-1})$ define an environment state and models for the behaviours of the other agents (= lower trees), but this is the only model they have
 - ❖ no “trying to figure out what the others must have done instead”
- ❖ so: maybe just “embrace the confusion” and assume the agent in the lower tree behaves randomly

Other Progress

- ❖ Performance was worrying, so I did some work to speed up the code
 - ❖ around 70-80% faster now
- ❖ Paper: wrote sections on I-POMDPs and modelling the system as an I-POMDP, starting to write about the planning algorithm
- ❖ Wrote lots of tests for code
- ❖ Will apply for time on Snellius today

Schedule

Goal: graduate on time (before Aug 31)

Effect of summer holidays?

Proposed schedule:

- ❖ simulations: beginning of July
- ❖ Meeting with Mike in July (after 17th)
- ❖ (Internal) paper draft deadline: e.g. July 24
- ❖ Submitting paper: 2 weeks before defence, e.g. Aug 4
- ❖ Defence: before Aug 23, e.g. Aug 18

Notes from the Meeting

- ❖ for particle noise, try both Gaussian and uniform distributions
- ❖ extend n_agents plots: 10 agents, also add bar plot
- ❖ investigate the lines “crossing over” in the n_agents plot
- ❖ exploration coefficient: choose value based on where convergence happens based on some objective. Maybe fix at 1.0, no need to do SA on it
- ❖ search depth: value between 4 and 6 (maybe preferably 4), need to do SA (up to 4)
- ❖ brand the mismatch issue as “bounded rationality” — it is natural that agents predict some action but something else happens instead
 - ❖ try to quantify the extent of mismatch: calculate frequencies in each tree
- ❖ make a better bar plot: multiple runs, different numbers of agents, confidence intervals
- ❖ estimate compute time for simulations, including how it scales with the number of agents
- ❖ in paper: make a simple diagram of the algorithm with two agents, two actions, draw out a few steps
- ❖ important to do global SA, will take time
- ❖ next meeting: July 19, noon-1pm
 - ❖ finish methods section before this, send Overleaf